

# Ridge And Lasso Regressions

*Jai Vrat Singh*

*23/07/2018*

Data used in this example can be downloaded from kaggle: <https://www.kaggle.com/c/boston-housing>

For myself, I have temporarily saved it into /Users/jvsingh/.kaggle/competitions/boston-housing using the kaggle api.

## Ridge Regression

```
filename="/Users/jvsingh/.kaggle/competitions/boston-housing/train.csv"
df = read.csv(file = filename, header = TRUE, stringsAsFactors = FALSE)
#uppercase columns
colnames(df) <- toupper(colnames(df))
head(df)
```

```
##      ID      CRIM      ZN  INDUS  CHAS      NOX      RM      AGE      DIS  RAD  TAX  PTRATIO
## 1  1 0.00632 18.0  2.31    0 0.538 6.575 65.2 4.0900    1 296    15.3
## 2  2 0.02731  0.0  7.07    0 0.469 6.421 78.9 4.9671    2 242    17.8
## 3  4 0.03237  0.0  2.18    0 0.458 6.998 45.8 6.0622    3 222    18.7
## 4  5 0.06905  0.0  2.18    0 0.458 7.147 54.2 6.0622    3 222    18.7
## 5  7 0.08829 12.5  7.87    0 0.524 6.012 66.6 5.5605    5 311    15.2
## 6 11 0.22489 12.5  7.87    0 0.524 6.377 94.3 6.3467    5 311    15.2
##      BLACK  LSTAT  MEDV
## 1 396.90   4.98 24.0
## 2 396.90   9.14 21.6
## 3 394.63   2.94 33.4
## 4 396.90   5.33 36.2
## 5 395.60  12.43 22.9
## 6 392.52  20.45 15.0
```

```
#we do not need ID
df[["ID"]] = NULL
head(df)
```

```
##      CRIM      ZN  INDUS  CHAS      NOX      RM      AGE      DIS  RAD  TAX  PTRATIO  BLACK
## 1 0.00632 18.0  2.31    0 0.538 6.575 65.2 4.0900    1 296    15.3 396.90
## 2 0.02731  0.0  7.07    0 0.469 6.421 78.9 4.9671    2 242    17.8 396.90
## 3 0.03237  0.0  2.18    0 0.458 6.998 45.8 6.0622    3 222    18.7 394.63
## 4 0.06905  0.0  2.18    0 0.458 7.147 54.2 6.0622    3 222    18.7 396.90
## 5 0.08829 12.5  7.87    0 0.524 6.012 66.6 5.5605    5 311    15.2 395.60
## 6 0.22489 12.5  7.87    0 0.524 6.377 94.3 6.3467    5 311    15.2 392.52
##      LSTAT  MEDV
## 1   4.98 24.0
## 2   9.14 21.6
## 3   2.94 33.4
## 4   5.33 36.2
## 5  12.43 22.9
## 6  20.45 15.0
```

```
X = df[,colnames(df)[-length(colnames(df))]]
Y = df[,length(colnames(df))]
```

## Ridge Model

```
#Ridge Model Fit
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 3.4.4
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## Loaded glmnet 2.0-16
```

```
ridge.mod = glmnet(x = as.matrix(X), y=Y, alpha=0, lambda = 0.1)
coef(ridge.mod)
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              s0
## (Intercept) 32.31334399
## CRIM        -0.049073407
## ZN          0.044489648
## INDUS       0.034325021
## CHAS        3.833756311
## NOX         -14.628904073
## RM          3.814889858
## AGE         -0.005497041
## DIS         -1.487442301
## RAD         0.288706424
## TAX         -0.010886780
## PTRATIO     -0.838062519
## BLACK       0.011615203
## LSTAT       -0.589050130
```

```
predict (ridge.mod, s=1, type = "coefficients")
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              1
## (Intercept) 32.31334399
## CRIM        -0.049073407
## ZN          0.044489648
## INDUS       0.034325021
## CHAS        3.833756311
## NOX         -14.628904073
## RM          3.814889858
## AGE         -0.005497041
## DIS         -1.487442301
## RAD         0.288706424
## TAX         -0.010886780
## PTRATIO     -0.838062519
## BLACK       0.011615203
## LSTAT       -0.589050130
```

```

#Prediction from API
predictions = predict (ridge.mod, s=1,newx = as.matrix(X))

#Prediction manually
thetas = as.matrix(predict (ridge.mod, s=1, type = "coefficients"))
predvals = cbind(rep(1,dim(X)[1]), as.matrix(X)) %*% thetas

#These values should be same
assertthat::are_equal(sum(abs(predictions - predvals)),0.0)

## [1] TRUE


```

Let us analyze errors

```

library(ggplot2)
# Histogram overlaid with kernel density curve
ggplot(data.frame(errors = as.numeric(Y - predictions)), aes(x=errors)) +
  geom_histogram(aes(y=..density..),      # Histogram with density instead of count on y-axis
                binwidth=.5,
                colour="black", fill="white") +
  geom_density(alpha=.2, fill="#FF6666") # Overlay with transparent density plot

```



```

#Let us produce ridge coefficients with different alphas
loglambdas = seq(from = -10, to = 20, length.out = 50)
lambdas     = exp(loglambdas)

RidgeCoeffs <- function(lambda){

```

```

ridge.mod = glmnet(x = as.matrix(X), y=Y, alpha=0, lambda = lambda)
as.matrix(coef(ridge.mod))
}

coeffsList = as.data.frame(t(sapply(lambdas, function(lambda) RidgeCoeffs(lambda))))

#prepare column names
colnames(coeffsList) = c("bias", colnames(df)[-length(colnames(df))])
coeffsDF = coeffsList
coeffsDF["loglambdas"] = loglambdas
head(coeffsDF)

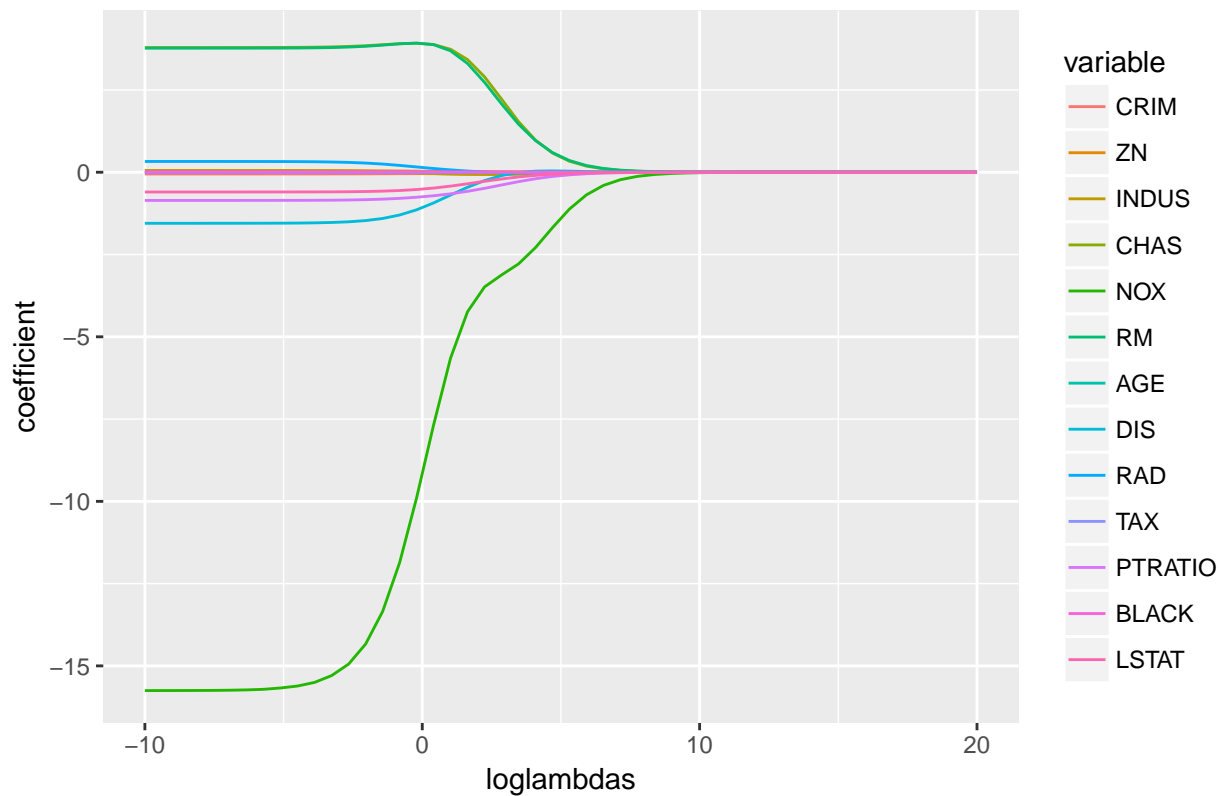
##      bias      CRIM      ZN      INDUS      CHAS      NOX      RM
## 1 34.03654 -0.05236777 0.04744169 0.05256458 3.788303 -15.74995 3.769750
## 2 34.03581 -0.05236625 0.04744038 0.05255684 3.788324 -15.74949 3.769772
## 3 34.03446 -0.05236344 0.04743797 0.05254256 3.788362 -15.74865 3.769812
## 4 34.03197 -0.05235827 0.04743352 0.05251622 3.788433 -15.74709 3.769886
## 5 34.02738 -0.05234874 0.04742533 0.05246768 3.788563 -15.74421 3.770023
## 6 34.01891 -0.05233118 0.04741022 0.05237822 3.788803 -15.73890 3.770275
##      AGE      DIS      RAD      TAX      PTRATIO      BLACK
## 1 -0.004666253 -1.551365 0.3274956 -0.01278266 -0.8561569 0.01165637
## 2 -0.004666682 -1.551339 0.3274787 -0.01278182 -0.8561497 0.01165636
## 3 -0.004667474 -1.551291 0.3274475 -0.01278027 -0.8561364 0.01165635
## 4 -0.004668935 -1.551203 0.3273898 -0.01277741 -0.8561119 0.01165632
## 5 -0.004671627 -1.551040 0.3272836 -0.01277213 -0.8560666 0.01165626
## 6 -0.004676587 -1.550740 0.3270878 -0.01276242 -0.8559832 0.01165615
##      LSTAT loglambdas
## 1 -0.6000606 -10.000000
## 2 -0.6000560 -9.387755
## 3 -0.6000476 -8.775510
## 4 -0.6000321 -8.163265
## 5 -0.6000034 -7.551020
## 6 -0.5999506 -6.938776

#plotting
library(reshape2)
coeffsDF[["bias"]] = NULL #Remove bias
coeffsDF.melt = melt(coeffsDF, id.vars = "loglambdas")
colnames(coeffsDF.melt) <- c("loglambdas", "variable", "coefficient")

ggplot(data=coeffsDF.melt, aes(x=loglambdas, y=coefficient, group=variable, colour=variable)) +
  geom_line() +
  ggtitle("Ridge Regression")

```

## Ridge Regression



## LASSO Regression

```
#Let us produce lasso coefficients with different alphas

LassoCoeffs <- function(lambda){
  mod = glmnet(x = as.matrix(X), y=Y, alpha=1, lambda = lambda)
  as.matrix(coef(mod))
}

coeffsList = as.data.frame(t(sapply(lambdas, function(lambda) LassoCoeffs(lambda))))

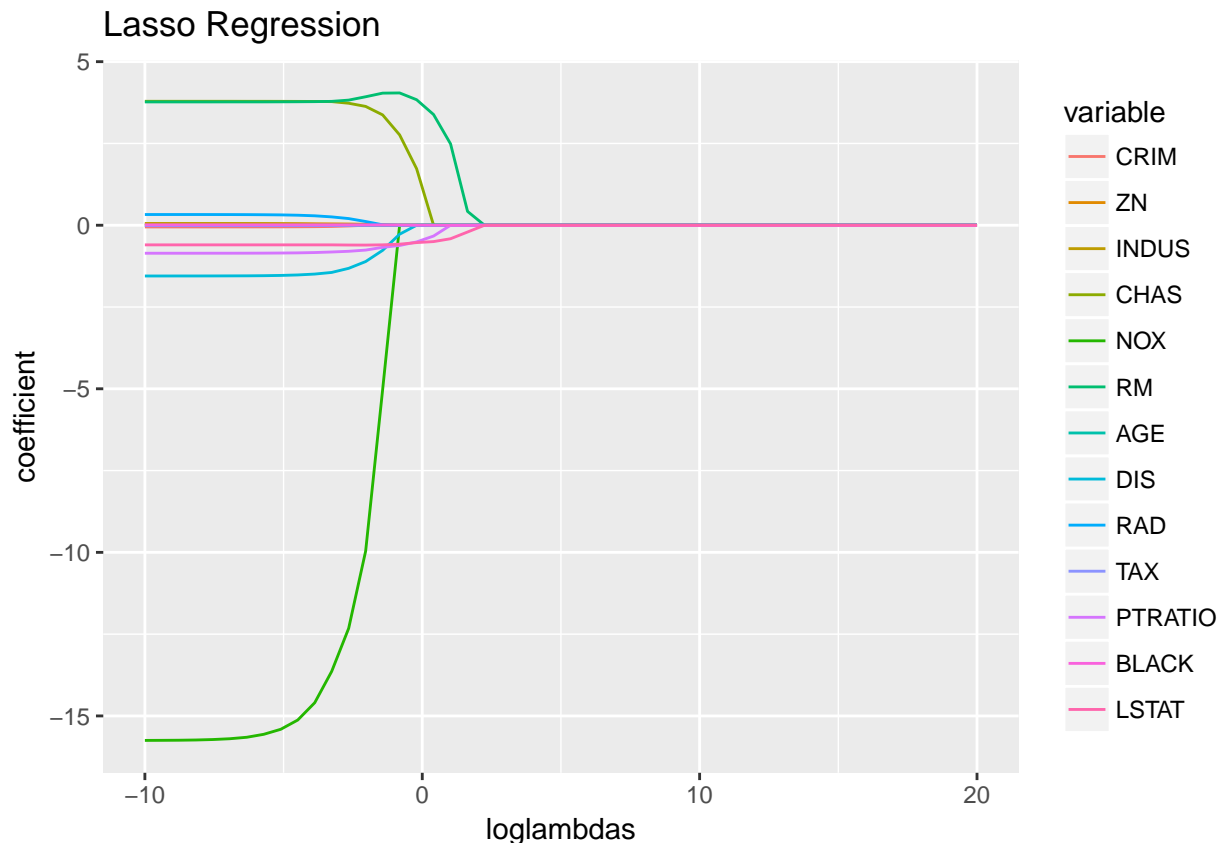
#prepare column names
colnames(coeffsList) = c("bias", colnames(df)[-length(colnames(df))])
coeffsDF = coeffsList
coeffsDF["loglambdas"] = loglambdas
head(coeffsDF)
```

```
##      bias      CRIM      ZN      INDUS      CHAS      NOX      RM
## 1 34.03425 -0.05235184 0.04743584 0.05251643 3.788266 -15.74798 3.769748
## 2 34.03158 -0.05233686 0.04742960 0.05246803 3.788256 -15.74586 3.769767
## 3 34.02665 -0.05230924 0.04741808 0.05237873 3.788238 -15.74195 3.769804
## 4 34.01757 -0.05225829 0.04739683 0.05221403 3.788203 -15.73474 3.769871
## 5 34.00081 -0.05216432 0.04735764 0.05191022 3.788140 -15.72143 3.769996
## 6 33.96990 -0.05199097 0.04728536 0.05134982 3.788024 -15.69689 3.770225
##      AGE      DIS      RAD      TAX      PTRATIO      BLACK
## 1 -0.004662889 -1.551265 0.3274294 -0.01277939 -0.8561208 0.01165602
## 2 -0.004660478 -1.551155 0.3273564 -0.01277578 -0.8560831 0.01165570
```

```
## 3 -0.004656031 -1.550952 0.3272219 -0.01276913 -0.8560135 0.01165513
## 4 -0.004647827 -1.550577 0.3269738 -0.01275686 -0.8558852 0.01165406
## 5 -0.004632694 -1.549887 0.3265160 -0.01273423 -0.8556485 0.01165211
## 6 -0.004604782 -1.548612 0.3256718 -0.01269248 -0.8552119 0.01164849
##      LSTAT loglambdas
## 1 -0.6000661 -10.000000
## 2 -0.6000662 -9.387755
## 3 -0.6000664 -8.775510
## 4 -0.6000667 -8.163265
## 5 -0.6000674 -7.551020
## 6 -0.6000685 -6.938776
```

```
#plotting
coeffsDF[["bias"]] = NULL #Remove bias
coeffsDF.melt = melt(coeffsDF, id.vars = "loglambdas")
colnames(coeffsDF.melt) <- c("loglambdas", "variable", "coefficient")

ggplot(data=coeffsDF.melt, aes(x=loglambdas, y=coefficient, group=variable, colour=variable)) +
  geom_line() +
  ggtitle("Lasso Regression")
```



As we see above, the Lasso converges faster to zero as expected with increasing regularization penalty

As seen below, both ridge and lasso have same estimates and equal to that of Linear Regression if we use the regularization parameter 0, i.e.  $\lambda=0$

```
ridge.mod = glmnet(x = as.matrix(X), y=Y, alpha=0, lambda = 0)
#as.matrix(t(coef(ridge.mod)))
lasso.mod = glmnet(x = as.matrix(X), y=Y, alpha=1, lambda = 0)
```

```

#as.matrix(t(coef(lasso.mod)))
reg.mod <- lm(MEDV ~ ., data = df)
#coef(reg.mod)
rbind(as.matrix(t(coef(ridge.mod))), as.matrix(t(coef(lasso.mod))), coef(reg.mod))

```

```

##      (Intercept)          CRIM          ZN          INDUS          CHAS          NOX
## s0      34.03741 -0.05236957 0.04744323 0.05257375 3.788278 -15.75050
## s0      34.03741 -0.05236957 0.04744323 0.05257375 3.788278 -15.75050
##      34.04544 -0.05248934 0.04744487 0.05385524 3.784864 -15.73966
##      RM          AGE          DIS          RAD          TAX          PTRATIO
## s0 3.769724 -0.004665744 -1.551396 0.3275157 -0.01278366 -0.8561655
## s0 3.769724 -0.004665744 -1.551396 0.3275157 -0.01278366 -0.8561655
##      3.768832 -0.004626602 -1.548823 0.3289671 -0.01286650 -0.8569757
##      BLACK          LSTAT
## s0 0.01165639 -0.6000660
## s0 0.01165639 -0.6000660
##      0.01166590 -0.6003155

```