

Volatility Modeling - ARCH/GARCH

Jai Vrat Singh

11/09/2018

R Markdown

```
ret.what <- list(date=numeric(), returns=numeric())
ret.widths <- c(12, 8)
strip.white <- c(TRUE, TRUE)

datF = scan(file = "http://faculty.chicagobooth.edu/ruey.tsay/teaching/fts2/m-3m4603.txt",
            what = ret.what, strip.white = strip.white)
datF = as.data.frame(datF)


#Above are simple returns, Convert returns to logReturns
datF[["logRet"]] = log(1+ datF$returns)

head(datF)
```

```
##      date  returns      logRet
## 1 19460228 -0.07792  0.075033258
## 2 19460330  0.01859 -0.018764966
## 3 19460430 -0.10000  0.095310180
## 4 19460531  0.20988 -0.235570446
## 5 19460628  0.00513 -0.005143204
## 6 19460731  0.07653 -0.079616965
```

Evidence of ARCH effect


```
plot(datF$logRet, type = "o", pch = ".")
```



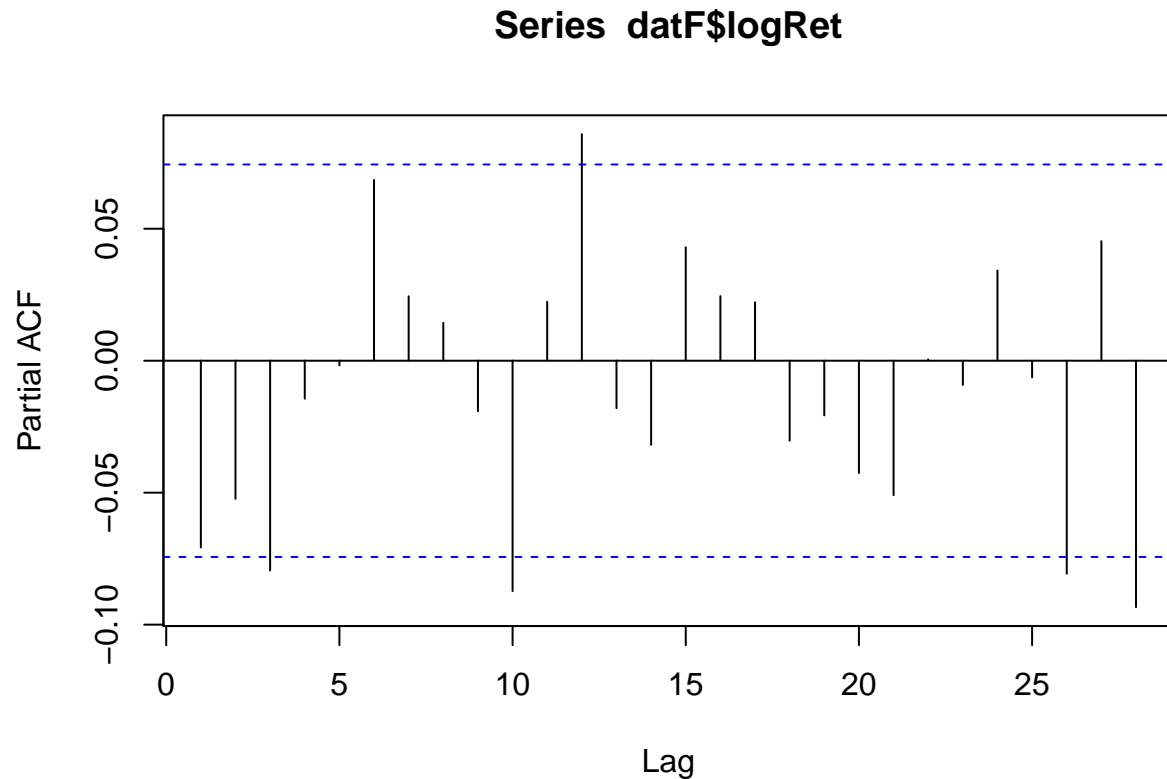
This plot reflects a pattern around the line $\log\text{Ret} = 0$. We need to look into ACF and PACF to get lagged correlations and its exact nature before the ARCH effect can be detected.

```
#par(mfrow=c(2,1))
acf(datF$logRet,)
```

Series datF\$logRet



```
acf(datF$logRet, type = "partial")
```



ACF till lags 12 => MA(12) PACF till 3, 12 => AR(3), AR(12)

Try these

```
#AR(12)
logret.fit = arima(x = datF$logRet, order = c(12,0,0))
logret.fit

##
## Call:
## arima(x = datF$logRet, order = c(12, 0, 0))
##
## Coefficients:
##      ar1      ar2      ar3      ar4      ar5      ar6      ar7      ar8
## -0.0838 -0.0500 -0.0708 -0.0027  0.0041  0.0649  0.0181  0.0088
## s.e.   0.0378   0.0379   0.0380   0.0383   0.0383   0.0383   0.0383   0.0383
##      ar9      ar10     ar11     ar12  intercept
## -0.0197 -0.0816  0.0307  0.0882    -0.0161
## s.e.   0.0383   0.0382   0.0384   0.0383     0.0023
##
## sigma^2 estimated as 0.004346:  log likelihood = 903.57,  aic = -1779.14
#aic = -1779.14

#AR(3)
logret.fit = arima(x = datF$logRet, order = c(3,0,0))
logret.fit
```

```
##
```

```
## Call:
## arima(x = datF$logRet, order = c(3, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3  intercept
##      -0.0794  -0.0584  -0.0808   -0.0161
## s.e.   0.0379   0.0379   0.0380    0.0021
##
## sigma^2 estimated as 0.004445:  log likelihood = 895.85,  aic = -1781.7
```

```
#aic = -1781.7
```

```
#MA(12)
```

```
logret.fit = arima(x = datF$logRet, order = c(0,0,12))
logret.fit
```

```
##
## Call:
## arima(x = datF$logRet, order = c(0, 0, 12))
##
## Coefficients:
##          ma1          ma2          ma3          ma4          ma5          ma6          ma7          ma8
##      -0.0821  -0.0402  -0.0691   0.0066   0.0121   0.0800   0.0165  -0.0001
## s.e.   0.0378   0.0378   0.0377   0.0382   0.0383   0.0386   0.0380   0.0390
##          ma9          ma10          ma11          ma12  intercept
##      -0.0316  -0.0909   0.0377   0.1035   -0.0161
## s.e.   0.0405   0.0396   0.0382   0.0376    0.0024
##
## sigma^2 estimated as 0.004338:  log likelihood = 904.24,  aic = -1780.48
```


```
#aic = -1780.48
```

Lowest AIC is for AR(3) => seems it is best fit

Residue analysis of above fit


```
#logret.fit$residuals
qqnormPlot(logret.fit$residuals)
```

NORM QQ PLOT




```
acf(logret.fit$residuals, lag.max = 20)
```

Series logret.fit\$residuals



```
acf(logret.fit$residuals, type = "partial", lag.max = 20)
```

Series logret.fit\$residuals



Trying various models and Residual analysis, seems that MA(12) is the best fit model.

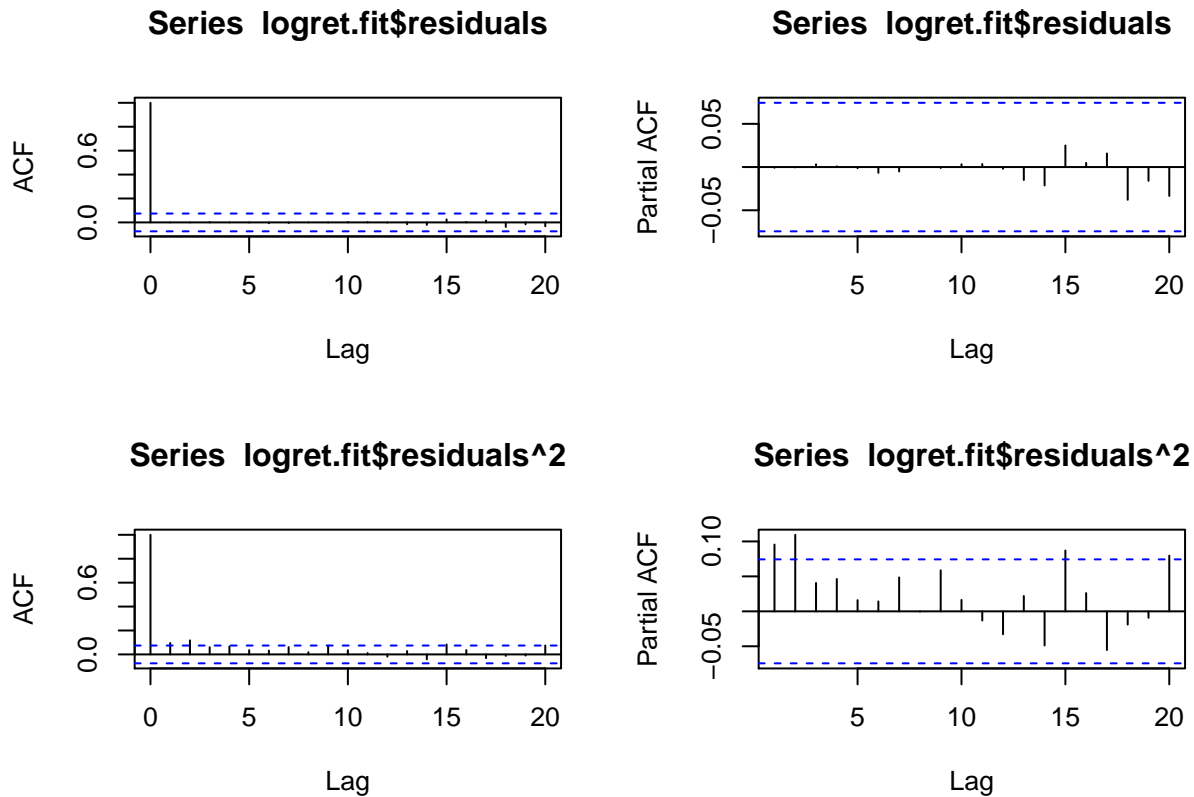
```
lags <- c(1:20)
sapply(lags, FUN = function(lag) {
  test <- Box.test(logret.fit$residuals, lag = lag, type = c("Box-Pierce", "Ljung-Box"))
  test$p.value
})
```

```
## [1] 0.9833876 0.9997333 0.9998227 0.9999916 0.9999995 0.9999986 0.9999996
## [8] 1.0000000 1.0000000 1.0000000 1.0000000 1.0000000 1.0000000 1.0000000
## [15] 0.9999998 0.9999999 1.0000000 0.9999980 0.9999988 0.9999952
```

P-values are very high => the residual is white noise.

ARCH effect on residuals

```
par(mfrow=c(2,2))
acf(logret.fit$residuals, lag.max = 20)
pacf(logret.fit$residuals, lag.max = 20)
acf(logret.fit$residuals^2, lag.max = 20)
pacf(logret.fit$residuals^2, lag.max = 20)
```



ACF and PACF of residuals clearly show conditional heteroscedasticity

Let us do LB test on residuals²

```
lags <- c(1:20)
sapply(lags, FUN = function(lag) {
  test <- Box.test(logret.fit$residuals^2, lag = lag, type = c("Box-Pierce", "Ljung-Box"))
  test$p.value
})
```


```
## [1] 0.0117694631 0.0003430990 0.0003504318 0.0002464178 0.0004326448
## [6] 0.0007452725 0.0005524597 0.0010513383 0.0004738406 0.0006691216
## [11] 0.0012090605 0.0019856451 0.0028282093 0.0031847927 0.0010630363
## [16] 0.0013212098 0.0017427494 0.0026919060 0.0041035567 0.0019481185
```

P-values are less than 0.05 => means that shocks are not independent. There is ARCH effect in almost all the lags.

PACF of squared log returns

```
pacf(datF$logRet^2, lag.max = 20)
```

Series datF\$logRet^2



Seems and ARCH(2) model can be fitted

$$r_t = \mu + a_t$$

$$a_t = \sigma_t \epsilon_t$$

$$\sigma_t^2 = \alpha_0 + \alpha_1 a_{t-1}^2 + \alpha_2 a_{t-2}^2$$

Fit the Model

```
library(fGarch)
garch2.fit <- garchFit(~garch(2,0), data = datF$logRet, trace = FALSE)
garch2.fit
```

```
##
## Title:
##  GARCH Modelling
##
## Call:
##  garchFit(formula = ~garch(2, 0), data = datF$logRet, trace = FALSE)
##
## Mean and Variance Equation:
##  data ~ garch(2, 0)
## <environment: 0x7fdf8c991580>
## [data = datF$logRet]
##
## Conditional Distribution:
##  norm
##
## Coefficient(s):
##      mu      omega    alpha1    alpha2
## -0.0169129  0.0035752  0.1125097  0.0985089
```



```
##
## Std. Errors:
## based on Hessian
##
## Error Analysis:
##      Estimate Std. Error t value Pr(>|t|)
## mu      -0.0169129  0.0024141  -7.006 2.46e-12 ***
## omega    0.0035752  0.0003063  11.672 < 2e-16 ***
## alpha1   0.1125097  0.0533135   2.110  0.0348 *
## alpha2   0.0985089  0.0511087   1.927  0.0539 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log Likelihood:
## 898.8208      normalized: 1.293267
##
## Description:
## Tue Sep 11 23:47:21 2018 by user:
```

$$\mu = -0.0169129$$

$$\alpha_0 = \omega = 0.0035752$$

$$\alpha_1 = 0.1125097$$

$$\alpha_2 = 0.0985089$$

So the Model is:


$$r_t = -0.0169129 + a_t$$

$$\sigma_t^2 = 0.0035752 + 0.1125097a_{t-1}^2 + 0.0985089a_{t-2}^2$$


$$a_t = \sigma_t \epsilon_t$$

```
resid <- residuals(garch2.fit)
par(mfrow=c(2,1))
acf(resid, lag.max = 20)
pacf(resid, lag.max = 20)
```

Series resid




Series resid



```
qqnormPlot(resid)
```

NORM QQ PLOT

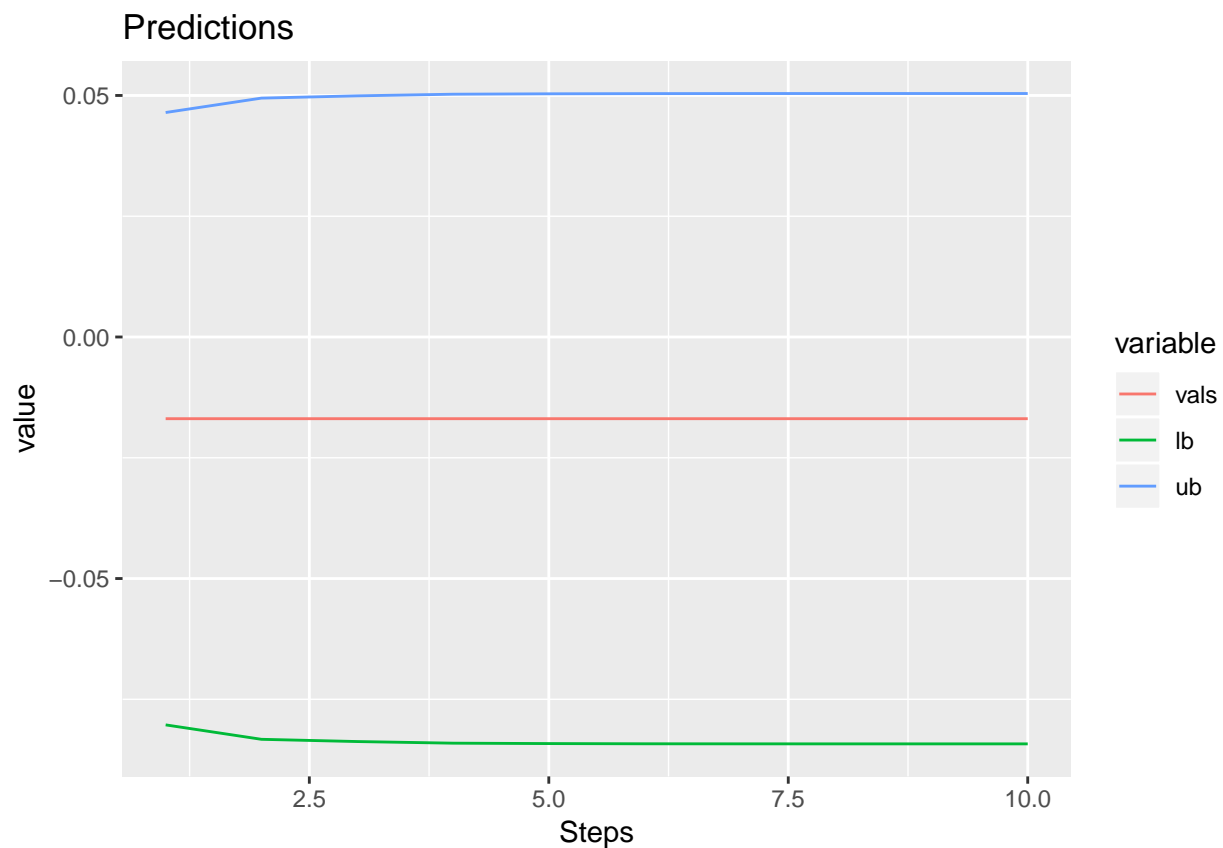


Prediction

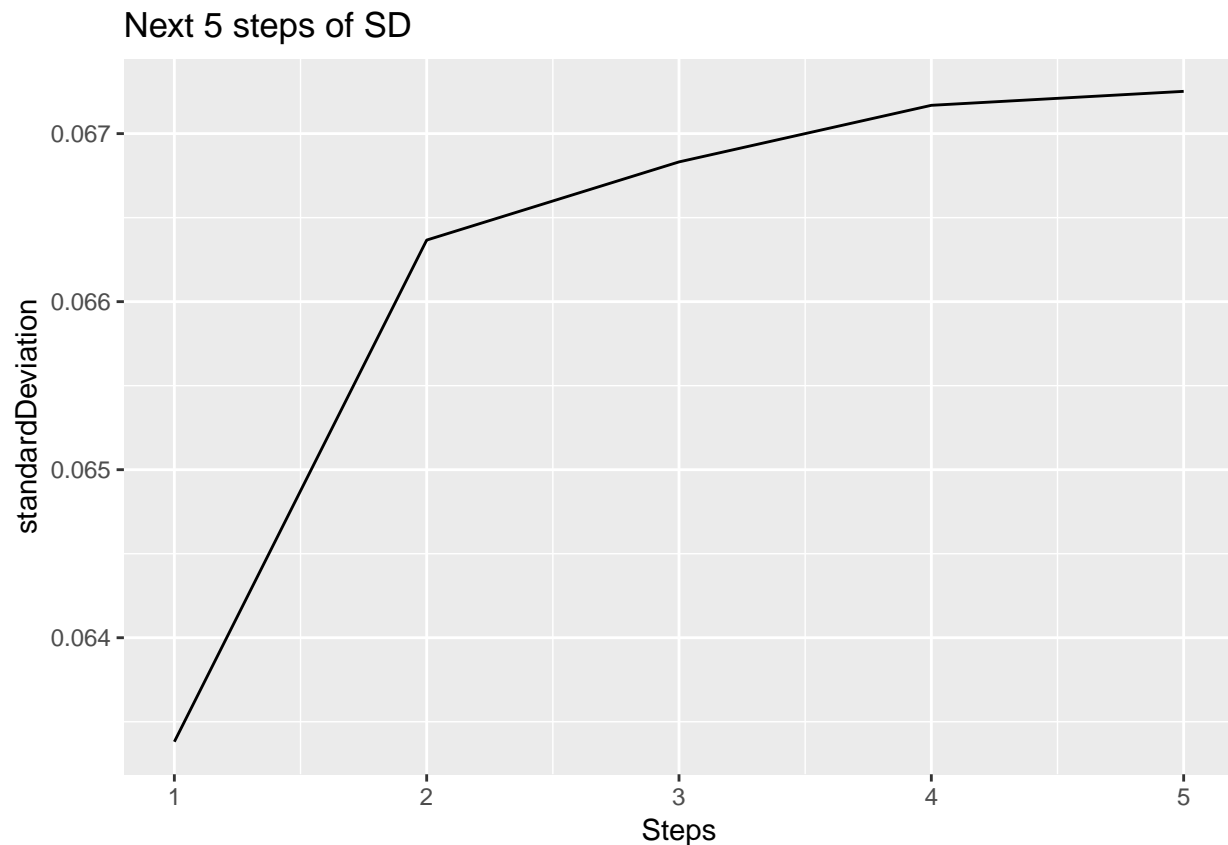
```
#predic.val <- predict(garch2.fit, 5)
predVal = fGarch::predict(garch2.fit, 10)

temp.dat <- data.frame("steps" = c(1:dim(predVal)[1]), vals = predVal$meanForecast,
                      lb = predVal$meanForecast - predVal$standardDeviation,
                      ub = predVal$meanForecast + predVal$standardDeviation)
temp.dat.melt = melt(temp.dat, id.vars = "steps")

ggplot(data=temp.dat.melt, aes(x=steps, y=value, group=variable, colour=variable)) +
  geom_line() +
  ggtitle("Predictions") +
  xlab("Steps")
```



```
#Predict Next 5 steps of SD
ggplot(data=predVal[1:5,], aes(x = c(1:dim(predVal)[1])[1:5], y=standardDeviation)) +
  geom_line() +
  ggtitle("Next 5 steps of SD") +
  xlab("Steps")
```



Bulding ARCH/GARCH model on S&P composite index

```
ret.what <- list(date=numeric(), gm=numeric(), sp=numeric())
ret.widths <- c(12, 25-13, 34-25)
strip.white <- c(TRUE, TRUE, TRUE)


datSP = scan(file = "./data/d-gmsp9303.txt", what = ret.what, strip.white = strip.white)
datSP = as.data.frame(datSP)

head(datSP)
```

```
##      date      gm      sp
## 1 19930104 0.01938 -0.000757
## 2 19930105 0.01141 -0.002389
## 3 19930106 0.02256 0.000414
## 4 19930107 -0.01838 -0.008722
## 5 19930108 0.00000 -0.003900
## 6 19930111 0.02622 0.004428
```


Graphically look for ARCH effect

```
plot(datSP$sp, type="o", pch=".")
```




```
acf(datSP$sp)
```

Series datSP\$sp




```
pacf(datSP$sp)
```

Series datSP\$sp




```
acf(datSP$sp^2)
```

Series datSP\$sp^2



```
pacf(datSP$sp^2)
```

Series datSP\$sp^2



The ACF plot shows almost no serial correlations but the squared series show significant autocorrelations.

Let us do auto correlations tests on few lags.

```
lags <- c(1:20)
sapply(lags, FUN = function(lag) {
  test <- Box.test(datSP$sp^2, lag = lag, type = "Ljung-Box", fitdf = 0)
  test$p.value
})
```

```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

The p-values are almost 0, this means we can reject the Null hypothesis that series are uncorrelated or independent.

Since serial correlations are present we need to fit the mean equation first and then apply ARCH test of residues.

```
#ARMA
arma6.6.fit = arima(x = datSP$sp, order = c(6,0,6))
arma6.6.fit
```

```
##
```

```
## Call:
```

```
## arima(x = datSP$sp, order = c(6, 0, 6))
```

```
##
```

```
## Coefficients:
```

```
## Warning in sqrt(diag(x$var.coef)): NaNs produced
```

```
##          ar1      ar2      ar3      ar4      ar5      ar6      ma1      ma2
##      0.3976  0.4411 -0.7451  0.5780 -0.0107 -0.6592 -0.4127 -0.459
## s.e.      NaN      NaN   0.1770  0.1557      NaN      NaN      NaN      NaN
```

```

##           ma3      ma4      ma5      ma6  intercept
##           0.7362 -0.5566 -0.0285 0.6448      4e-04
## s.e.    0.2148  0.1533  0.1008    NaN      2e-04
##
## sigma^2 estimated as 0.0001197:  log likelihood = 8583.27,  aic = -17138.54
arima5.5.fit = arima(x = datSP$sp, order = c(5,0,5))
arima5.5.fit

##
## Call:
## arima(x = datSP$sp, order = c(5, 0, 5))
##
## Coefficients:
##           ar1      ar2      ar3      ar4      ar5      ma1      ma2      ma3
##           0.6955  0.6782 -0.6275 -0.1178 -0.0894 -0.7125 -0.6990  0.6303
## s.e.    0.4136  0.8117  0.6045  0.6350  0.6036  0.4146  0.8152  0.6185
##           ma4      ma5  intercept
##           0.1591  0.0403      4e-04
## s.e.    0.6197  0.5935      2e-04
##
## sigma^2 estimated as 0.0001198:  log likelihood = 8581.18,  aic = -17138.35
#aic = -17138.35

arima4.4.fit = arima(x = datSP$sp, order = c(4,0,4))
arima4.4.fit

##
## Call:
## arima(x = datSP$sp, order = c(4, 0, 4))
##
## Coefficients:
##
## Warning in sqrt(diag(x$var.coef)): NaNs produced
##           ar1      ar2      ar3      ar4      ma1      ma2      ma3      ma4
##           -0.1362  0.2585  0.0087  0.3602  0.1160 -0.2906 -0.0536 -0.347
## s.e.    0.2767    NaN  0.1610    NaN  0.2735    NaN  0.1673    NaN
##           intercept
##           4e-04
## s.e.    2e-04
##
## sigma^2 estimated as 0.0001206:  log likelihood = 8572.39,  aic = -17124.79
#Fails with Nan

arima2.2.fit = arima(x = datSP$sp, order = c(2,0,2))
arima2.2.fit

##
## Call:
## arima(x = datSP$sp, order = c(2, 0, 2))
##
## Coefficients:
##           ar1      ar2      ma1      ma2  intercept
##           0.7499 -0.0148 -0.7646 -0.0114      4e-04
## s.e.    0.8820  0.7469  0.8735  0.7599      2e-04

```



```
##
## sigma^2 estimated as 0.0001208: log likelihood = 8569.87, aic = -17127.73
```

```
#aic = -17127.73
```

```
ar6.fit = arima(x = datSP$sp, order = c(6,0,0))
ar6.fit
```

```
##
## Call:
## arima(x = datSP$sp, order = c(6, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ar6  intercept
##      -0.0133  -0.0279  -0.0362   0.0033  -0.0502  -0.0204         4e-04
## s.e.   0.0190   0.0190   0.0190   0.0190   0.0190   0.0190         2e-04
##
## sigma^2 estimated as 0.0001206: log likelihood = 8572.24, aic = -17128.48
```

```
#aic = -17128.48
```

```
ar5.fit = arima(x = datSP$sp, order = c(5,0,0))
ar5.fit
```


```
##
## Call:
## arima(x = datSP$sp, order = c(5, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5  intercept
##      -0.0123  -0.028  -0.0355   0.0039  -0.050         4e-04
## s.e.   0.0190   0.019   0.0190   0.0190   0.019         2e-04
##
## sigma^2 estimated as 0.0001207: log likelihood = 8571.66, aic = -17129.32
```

```
#aic = -17129.32
```

We pick ARMA(5,5) model for lowest AIC values


```
arima5.5.fit = arima(x = datSP$sp, order = c(5,0,5))
```

```
lags <- c(1:20)
pvals <- sapply(lags, FUN = function(lag) {
  test <- Box.test(arima5.5.fit$residuals, lag = lag, type = "Ljung-Box", fitdf = 0)
  test$p.value
})
plot(x=lags, y = pvals)
```

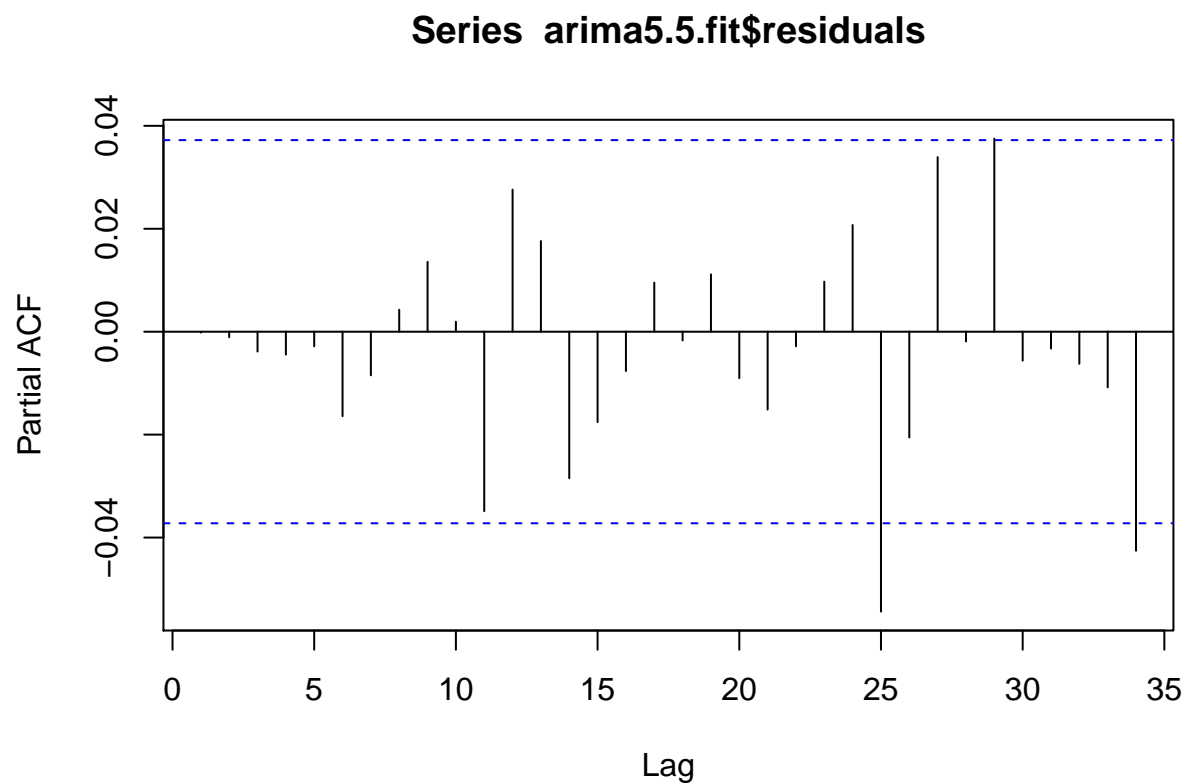


```
#ACF/PACF  
acf(arima5.5.fit$residuals)
```

Series arima5.5.fit\$residuals




```
pacf(arima5.5.fit$residuals)
```



So the residuals are not autocorrelated.


ARCH test on residual squared:

```
lags <- c(1:20)
pvals <- sapply(lags, FUN = function(lag) {
  test <- Box.test(arima5.5.fit$residuals^2, lag = lag, type = "Ljung-Box", fitdf = 0)
  test$p.value
})
plot(x=lags, y = pvals)
```




```
#ACF/PACF  
acf(arima5.5.fit$residuals^2)
```

Series arima5.5.fit\$residuals^2



```
pacf(arima5.5.fit$residuals^2)
```

Series `arma5.5.fit$residuals^2`



So there is autocorrelation in residues.

To build a garch model , we do a joint max likelihood estimation with AR(5). Not that AARMA(5,5) does not converge/gives NaNs

```
#garch55.fit <- garchFit(~garch(5,5), data = datSP$sp, trace = FALSE)
#garch55.fit
```

```
library(rugarch)
```

```
## Loading required package: parallel
```

```
##
```

```
## Attaching package: 'rugarch'
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
##      sigma
```

```
spec <- ugarchspec(variance.model = list(model = "sGARCH",
                                          garchOrder = c(1, 1),
                                          submodel = NULL,
                                          external.regressors = NULL,
                                          variance.targeting = FALSE),

                  mean.model      = list(armaOrder = c(5, 0),
                                          external.regressors = NULL,
                                          distribution.model = "norm",
                                          start.pars = list(),
                                          fixed.pars = list()))
```

```
## Warning: unidentified option(s) in mean.model:
```

```
## distribution.model start.pars fixed.pars
garch.ar5.fit <- ugarchfit(spec = spec, data = datSP$sp, solver.control = list(trace=0))

#garch.ar5.fit <- garchFit(formula = ~garch(2,1), data = datSP$sp, trace = FALSE)
garch.ar5.fit

##
## *-----*
## *          GARCH Model Fit          *
## *-----*
##
## Conditional Variance Dynamics
## -----
## GARCH Model   : sGARCH(1,1)
## Mean Model    : ARFIMA(5,0,0)
## Distribution   : norm
##
## Optimal Parameters
## -----
##      Estimate Std. Error t value Pr(>|t|)
## mu      0.000640  0.000134  4.78344 0.000002
## ar1      0.013345  0.019855  0.67214 0.501493
## ar2     -0.005858  0.019653 -0.29807 0.765648
## ar3     -0.039025  0.019652 -1.98588 0.047047
## ar4     -0.020464  0.019767 -1.03524 0.300556
## ar5     -0.072067  0.019518 -3.69233 0.000222
## omega    0.000001  0.000001  0.89343 0.371630
## alpha1   0.069731  0.010754  6.48415 0.000000
## beta1    0.928146  0.010498 88.41024 0.000000
##
## Robust Standard Errors:
##      Estimate Std. Error t value Pr(>|t|)
## mu      0.000640  0.000190  3.36860 0.000756
## ar1      0.013345  0.021742  0.61381 0.539339
## ar2     -0.005858  0.021299 -0.27504 0.783288
## ar3     -0.039025  0.019726 -1.97836 0.047888
## ar4     -0.020464  0.020303 -1.00795 0.313477
## ar5     -0.072067  0.018984 -3.79612 0.000147
## omega    0.000001  0.000004  0.13064 0.896060
## alpha1   0.069731  0.094012  0.74173 0.458253
## beta1    0.928146  0.088974 10.43170 0.000000
##
## LogLikelihood : 8977.043
##
## Information Criteria
## -----
##
## Akaike      -6.4704
## Bayes       -6.4512
## Shibata     -6.4705
## Hannan-Quinn -6.4635
##
## Weighted Ljung-Box Test on Standardized Residuals
## -----
```

```


##                                statistic p-value
## Lag[1]                        0.147  0.7014
## Lag[2*(p+q)+(p+q)-1][14]     5.249  1.0000
## Lag[4*(p+q)+(p+q)-1][24]     11.923  0.5598
## d.o.f=5
## H0 : No serial correlation
##
## Weighted Ljung-Box Test on Standardized Squared Residuals
## -----
##                                statistic p-value
## Lag[1]                        0.02539  0.8734
## Lag[2*(p+q)+(p+q)-1][5]      3.10970  0.3874
## Lag[4*(p+q)+(p+q)-1][9]      3.64911  0.6488
## d.o.f=2
##
## Weighted ARCH LM Tests
## -----
##          Statistic Shape Scale P-Value
## ARCH Lag[3]    0.07269 0.500 2.000  0.7875
## ARCH Lag[5]    0.17421 1.440 1.667  0.9715
## ARCH Lag[7]    0.26453 2.315 1.543  0.9945
##
## Nyblom stability test
## -----
## Joint Statistic:  457.439
## Individual Statistics:
## mu      0.18426
## ar1     0.67833
## ar2     0.06316
## ar3     0.12635
## ar4     0.32792
## ar5     0.02737
## omega   97.02699
## alpha1  0.29524
## beta1   0.23932
##
## Asymptotic Critical Values (10% 5% 1%)
## Joint Statistic:      2.1 2.32 2.82
## Individual Statistic:  0.35 0.47 0.75
##
## Sign Bias Test
## -----
##          t-value      prob sig
## Sign Bias      1.970 4.892e-02 **
## Negative Sign Bias  0.555 5.789e-01
## Positive Sign Bias  2.292 2.200e-02 **
## Joint Effect     29.492 1.765e-06 ***
##
##
## Adjusted Pearson Goodness-of-Fit Test:
## -----
## group statistic p-value(g-1)
## 1    20    55.85    1.721e-05
## 2    30    67.59    6.434e-05

```

```
## 3    40    90.25    6.099e-06
## 4    50    96.51    6.053e-05
##
##
## Elapsed time : 0.501334
```


All coefficients are significantly different from 0

```
resid <- residuals(garch.ar5.fit)
lags <- c(1:10)
pvals <- sapply(lags, FUN = function(lag) {
  test <- Box.test(resid, lag = lag, type = "Ljung-Box", fitdf = 0)
  test$p.value
})
plot(x=lags, y = pvals)
abline(h = 0.05)
```




```
#ACF/PACF
acf(resid, lag.max = 10)
```


Series resid



```
pacf(resid, lag.max = 10)
```

Series resid



```
predVal = ugarchforecast(garch.ar5.fit,n.ahead=10,data=datSP$sp)
```

```

#predVal@forecast$seriesFor
#predVal@forecast$sigmaFor

oldPlusForeCast <- c(datSP$sp, predVal@forecast$seriesFor)
ub               <- c(datSP$sp, predVal@forecast$seriesFor + predVal@forecast$sigmaFor)
lb               <- c(datSP$sp, predVal@forecast$seriesFor - predVal@forecast$sigmaFor)
temp.dat <- data.frame(x = 1:length(oldPlusForeCast) , vals = oldPlusForeCast, lb = lb, ub=ub)

#Last 30 + 10 projected means and theor bounds
temp.dat.melt = melt(tail(temp.dat,50), id.vars = "x")

ggplot(data=temp.dat.melt, aes(x=x, y=value, group=variable, colour=variable)) +
  geom_line()+
  ggtitle("Current + Pred")+
  xlab("x")

```

